

CS162  
Operating Systems and  
Systems Programming  
Lecture 2

Protection: Processes and Kernels

Professor Natacha Crooks & Matei Zaharia

<https://cs162.org/>



# Homework and Early Drop Deadline

---

HW0 Due 30/1

Early Drop Deadline 31/1

Should be working on Homework 0 already!

Get familiar with all the cs162 tools

HW1 will be released on 31/1

# Projects are looming

---

Group Formation Form (Link on EdStem) is due Feb 8th.

There is a teammate search functionality on EdStem.

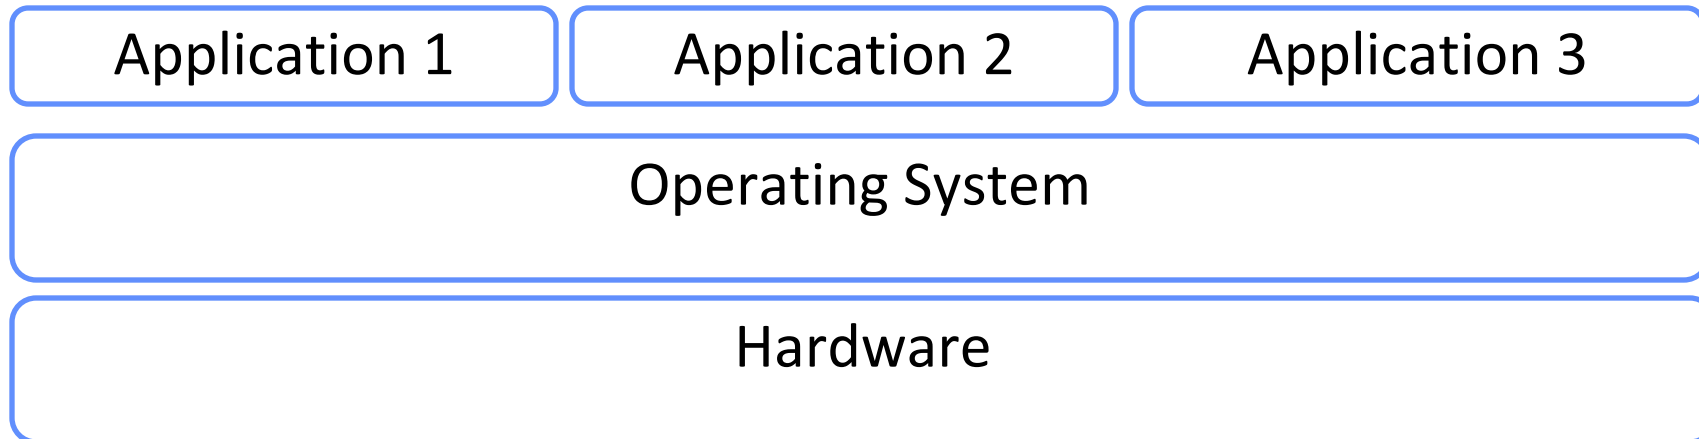
Discussions are starting! First 2 optional but mandatory afterwards

Project 0 will be released on 27/1

# Recall: Operating System

---

An operating system implements a **virtual machine** for the application whose interface is more **convenient** than the raw hardware interface  
(convenient = security, reliability, portability)



# Recall: Three main hats

---




## Referee

Manage protection, isolation, and sharing of resources



## Illusionist

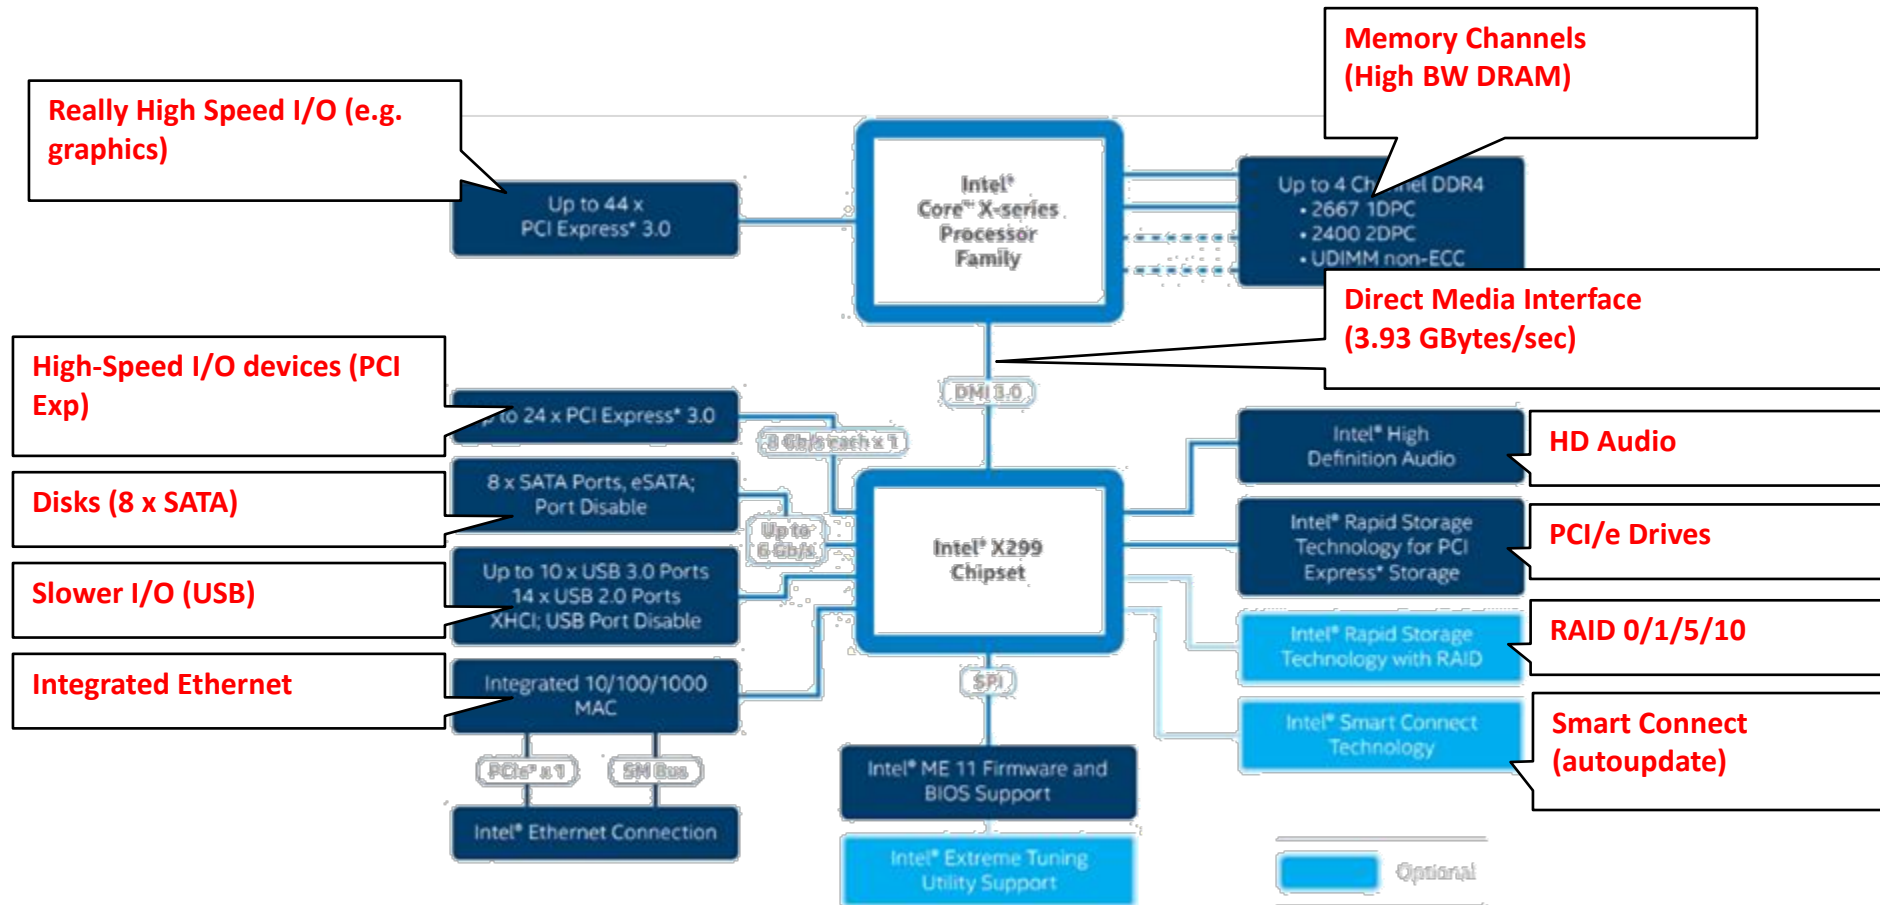
Provide clean, easy-to-use abstractions of physical resources



## Glue

Provides a set of common services


# Recall: HW Complex



## Intel Skylake-X I/O Configuration

# Recall: Increasing Software Complexity

---





# Evaluation Criteria

---

## **Performance**

Efficiently, with low overhead  
and equitably.

## **Reliability**

System does what it supposed to do

## **Security**

Minimise vulnerability to attack

## **Portability**

No need to change abstractions  
when hardware changes

# Evaluation Criteria: Reliability

---

System does what it is supposed to do

OS failures catastrophic!



**Availability:** mean time to failure + mean time to repair

# Evaluation Criteria: Security

Minimize vulnerability to attack

**Integrity:** Computer's operation cannot be compromised by a malicious attacker

**Privacy:** data stored on computer accessible to authorized users

## Enforcement Policy

How the OS ensures only permitted actions are allowed

## Security Policy

What is permitted




# Evaluation Criteria: Portability

---

A portable abstraction **does not change** as the hardware changes

Can't rewrite application (or OS!) every time

Must plan for hardware that does not exist yet!



# Three “Prongs” for the Class

---

Understanding OS principles

System Programming

Map Concepts to Real Code

# Topic Breakdown

---

Virtualizing the CPU

Process Abstraction and API

Threads and Concurrency

Scheduling

Virtualizing Memory

Virtual Memory

Paging

Persistence

IO devices

File Systems

Distributed Systems

Challenges with distribution

Data Processing & Storage

# Topic Breakdown

---

Virtualizing the CPU

Process Abstraction and API

Threads and Concurrency

Scheduling

Virtualizing Memory

Virtual Memory

Paging

Persistence

IO devices

File Systems

Distributed Systems

Challenges with distribution

Data Processing & Storage

## Side Note: Mechanisms vs Policy

---

### Mechanism

Low-level methods or protocols  
that implement a needed piece of  
functionality

**A Brake Pedal!**

### Policy

Algorithms for making decisions  
within the OS.  
Use the mechanism.

**“I break when I see a stop sign”**



# Goals for Today

---

- What are the requirements of a good VM abstraction?
- What is a **process**?
- How does the **kernel** use processes to enforce protection?
- When does one switch from **kernel** to **user mode** and back?

# Goal 1: Requirements for Virtualization

---

# The OS will protect you

---

Protection is necessary to preserve the virtualization abstraction

Protect applications from other application's code  
(reliability, security, privacy)

Protect OS from the application

Protect applications against inequitable resource utilisation  
(memory, CPU time)



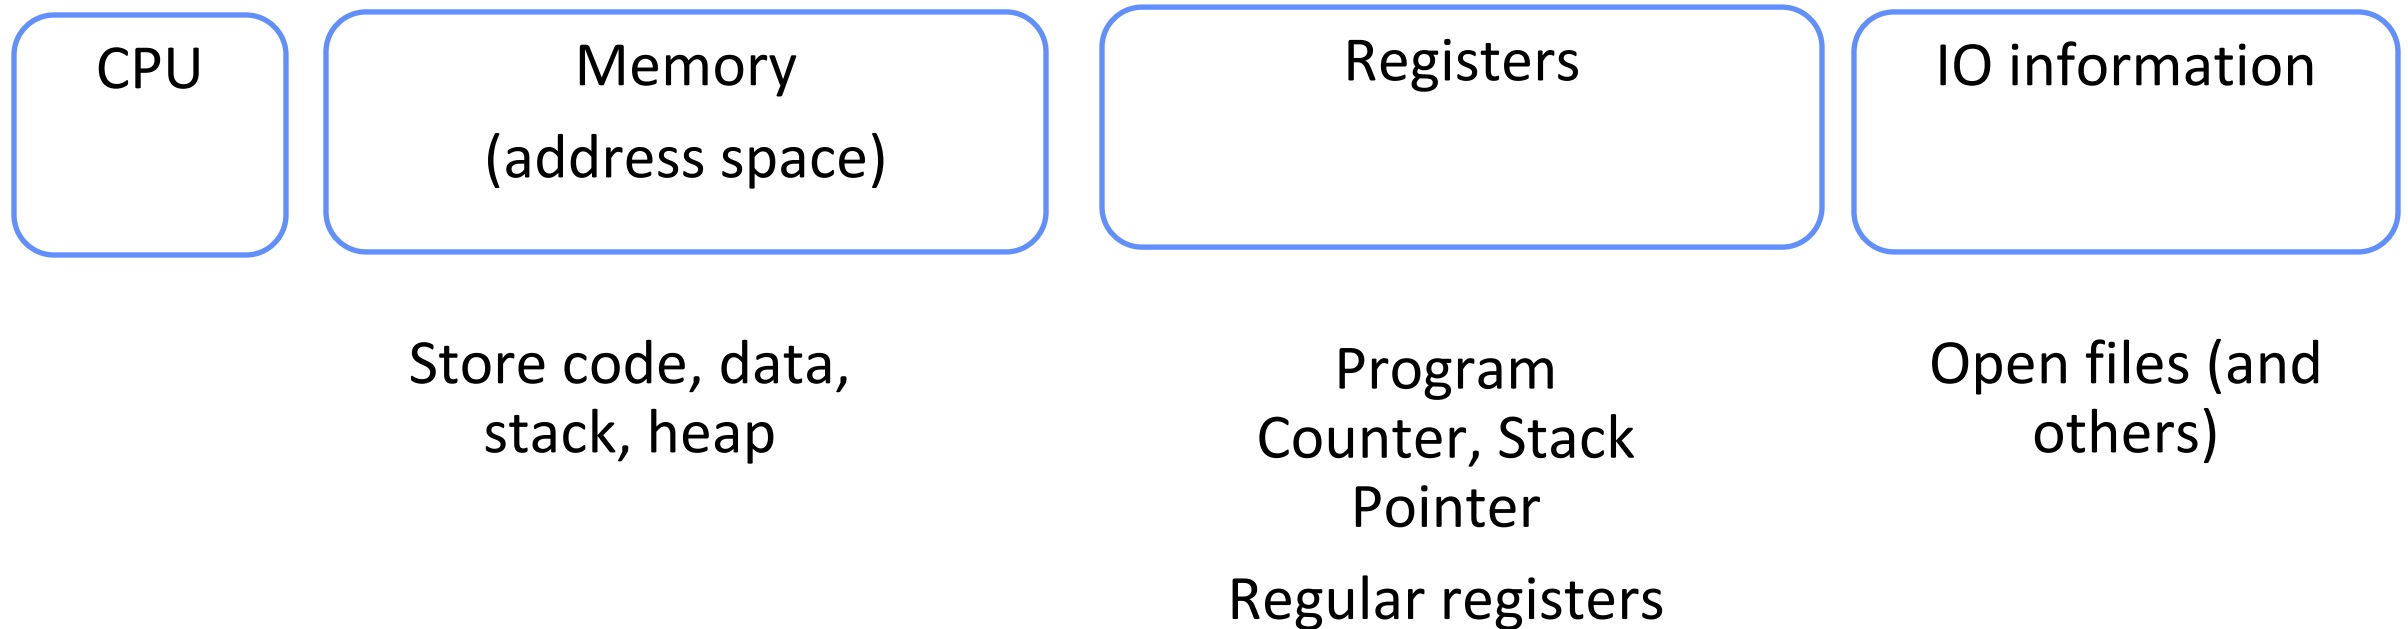
## Goal 2: What is a Process?

---


# A process (simplified)

---

A process is an **instance** of a running program




# From program to process




```
#include <stdio.h>

int main()
{
    printf("Hello World");
    return 0;
}
```





Executable image,  
instructions and  
data  
  
./helloworld

crooks@laptop> ./helloworld



crooks@laptop> ./helloworld



Task Manager

File Options View

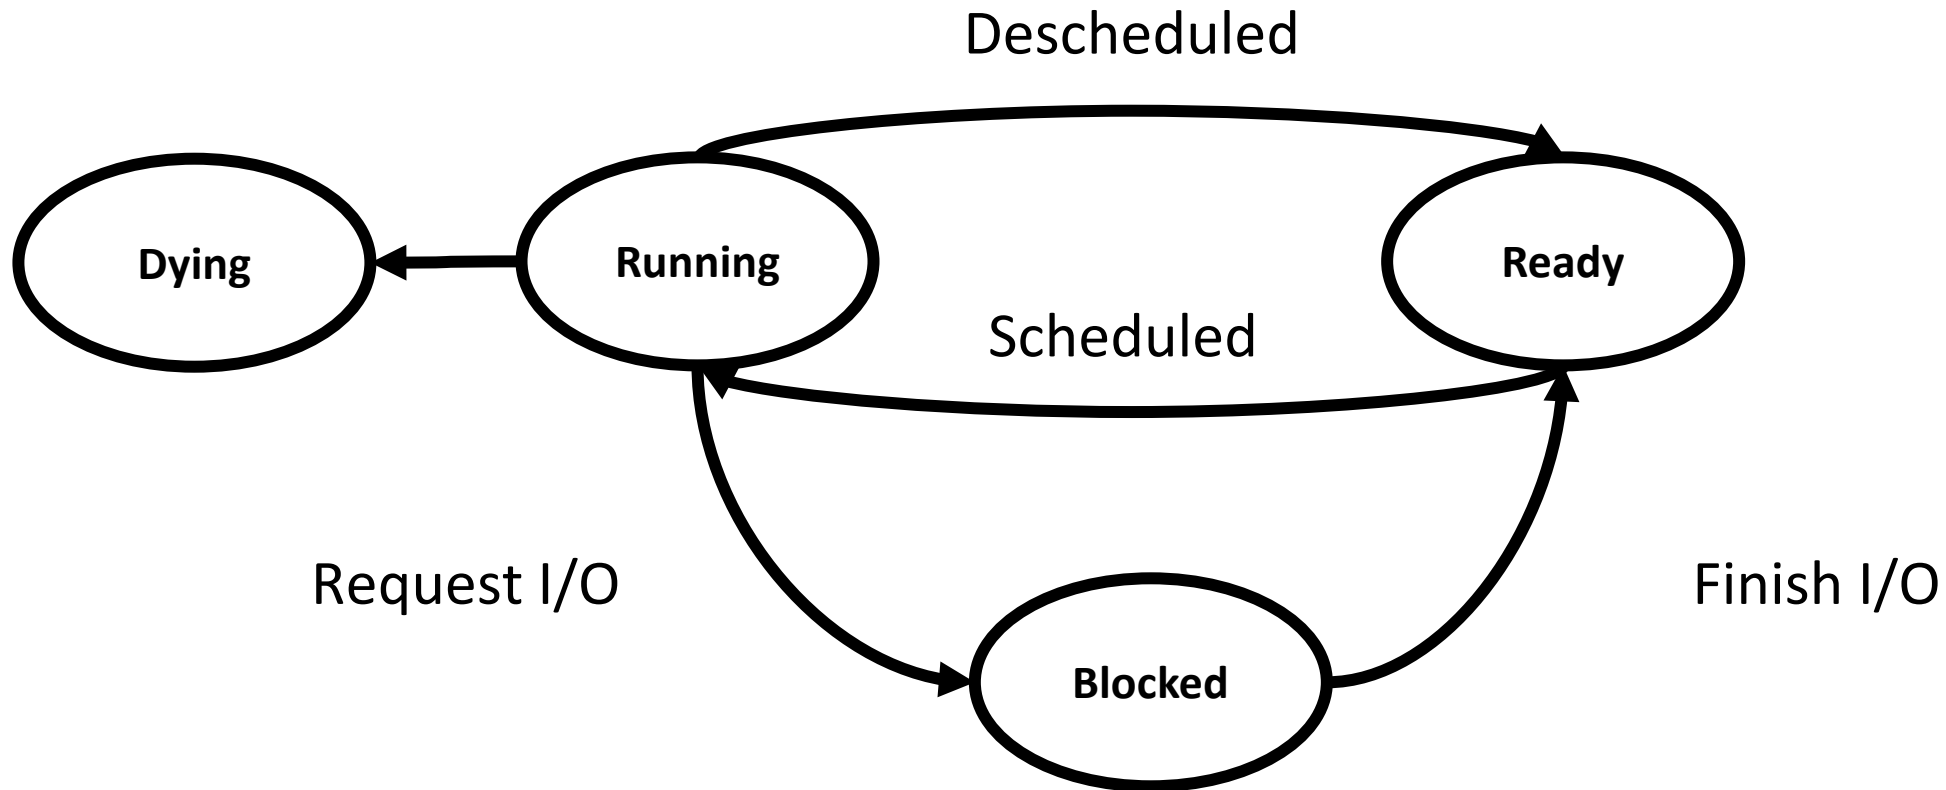
Processes Performance App history Startup Users Details Services

Name	Status	5% CPU	50% Memory	1% Disk	0% Network	2% GPU	GPU engine	Power usage	Power usage tr...
> Google Chrome (41)		0.9%	3,390.9 MB	0.1 MB/s	0.1 Mbps	1.7%	GPU 0 - Video Decode	Very low	Very low
> Microsoft PowerPoint (2)		0.1%	735.1 MB	0 MB/s	0.1 Mbps	0%		Very low	Very low
> Slack (6)		0%	413.2 MB	0 MB/s	0.1 Mbps	0%		Very low	Very low

# Process Life Cycle

---

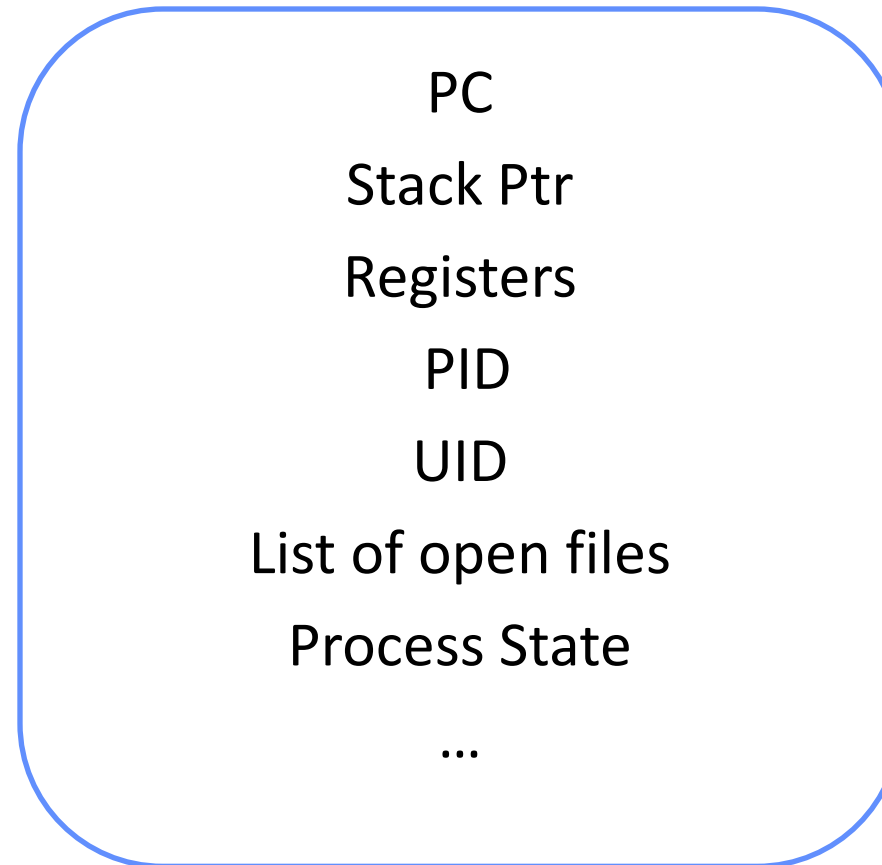
A process can be in one of several states:  
(real OSes have additional variants)



# Process Management by the OS

---

Process Control Block (or process descriptor)  
in OS stores necessary metadata





# Three “Prongs” for the Class

---

Understanding OS principles

System Programming

Map Concepts to Real Code

# Processes in the wild (well, in the kernel)

```
enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };

// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)
};
```

In Linux: `task_struct` defined  
in `<linux/sched.h>`

# Processes in Pintos

---

```
struct process {
    /* Owned by process.c. */
    uint32_t* pagedir;          /* Page directory. */
    char process_name[16];      /* Name of the main thread */
    struct thread* main_thread; /* Pointer to main thread */

    /* All the fun data structures you're going to add */
};
```

Pintos (userprog/process.h)

# Many Processes

---

Process List stores all processes

```
struct {  
    struct spinlock lock;  
    struct proc proc[NPROC];  
} ptable;
```

Xv6 Kernel (proc.c)

Run Queues

Lists all PCBs in **READY** state

Wait Queues

Lists all PCBs in **BLOCKED** state

# The Illusionist and the Referee are Back

---



Illusionist

Give every **process** the illusion of running on a private CPU

Give every **process** the illusion of running on private memory



Referee

Manage resources to allocate to each **process**

Isolate **process** from all other processes and protect OS



# Operating System Kernel

---

Lowest level of OS running on system.

Kernel is **trusted** with **full access** to all hardware capabilities

All other software (OS or applications) is considered **untrusted**

Untrusted

Applications

Rest of OS

Trusted

Operating System Kernel


Untrusted

Hardware

# The Process, Refined

---

A executing program with **restricted rights**



Enforcing mechanism must not hinder **functionality** or hurt **performance**

# User vs Kernel: Dr Jekyll and Mr Hyde

---



Application/User Code  
(Untrusted)

Run all the processor with all  
potentially dangerous  
operations disabled

Kernel Code (Trusted)

Runs directly on processor with  
unlimited rights

Performs any hardware  
operations

But run on the same machine!



# How can the kernel enforce restricted rights?

---

1) While preserving functionality

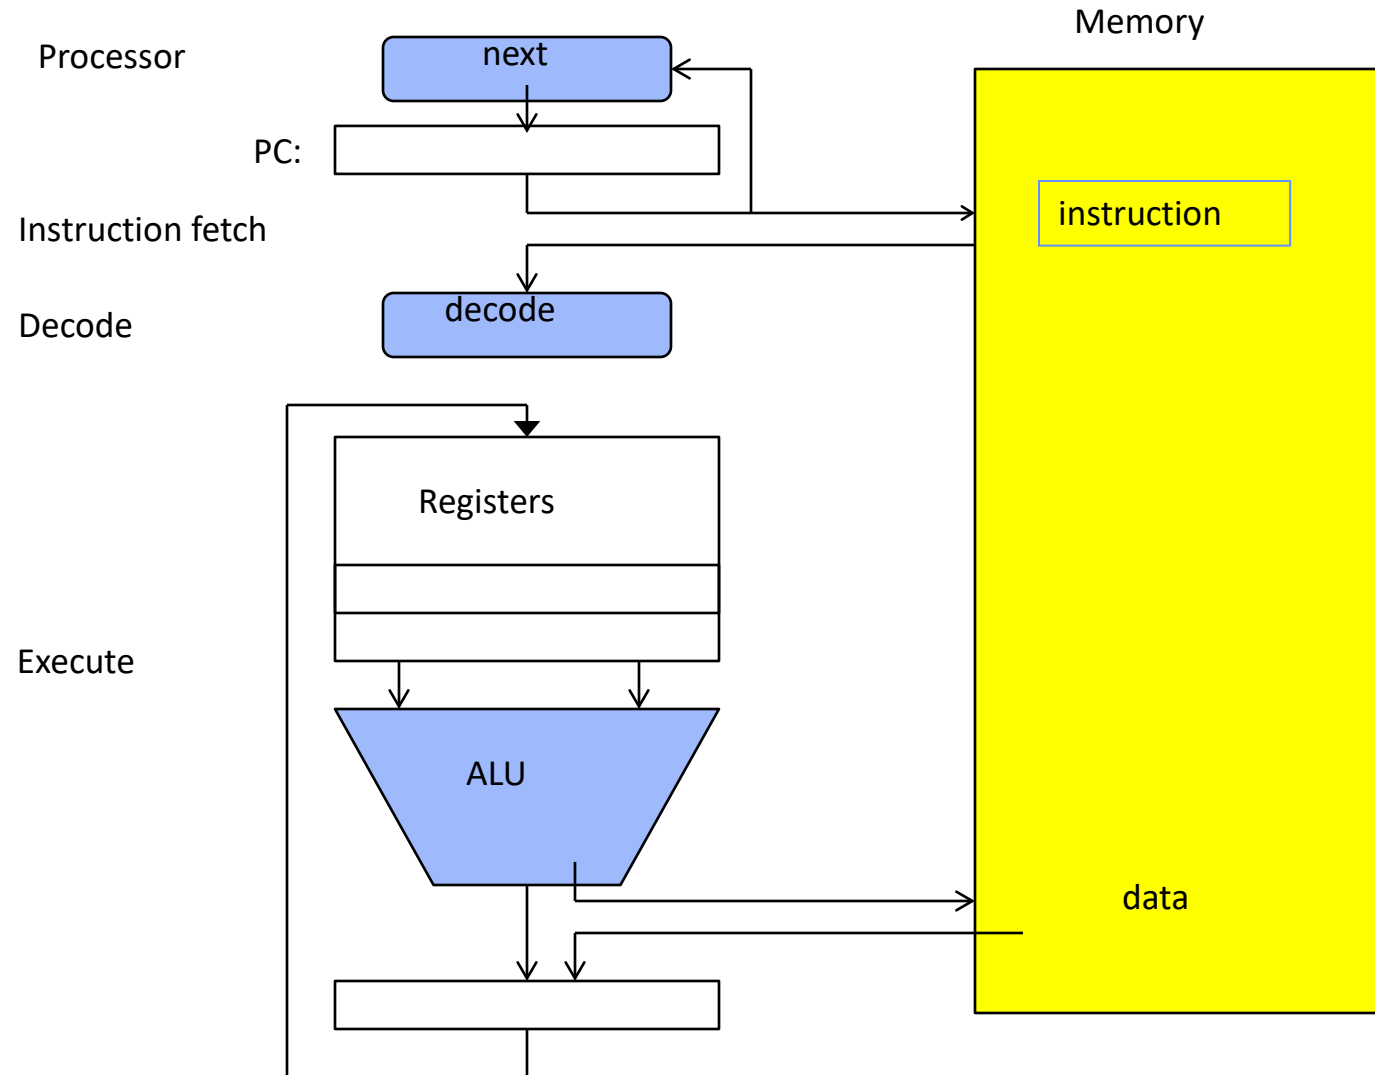
2) While preserving performance

3) While preserving control

# Attempt 1: Simulation


---

# Recall: CPU Instruction Cycle (from CS61c)



# Attempt 1: Simulation

---



## Potential Issues:

**Extremely slow!** Would have to cycle through all operations, switch into the kernel, etc.

**Unnecessary.** Most operations are perfectly safe!

# Attempt 2: Dual Mode Operation

---

Hardware to the rescue!

Use a bit to enable two modes of execution

## In User Mode

Processor checks each instruction before executing it

Executes a limited (safe) set of instructions

## In Kernel Mode

OS executes with protection checks off

Can execute any instructions



# Hardware must support

---

## 1) Privileged Instructions

Unsafe instructions cannot be executed in user mode

## 2) Memory Isolation

Memory accesses outside a process's address space prohibited

## 3) Interrupts

Ensure kernel can regain control from running process

## 4) Safe Transfers

Correctly transfer control from user-mode to kernel-mode and back

## Req 1/4: Privileged Instructions

---

Cannot change privilege level (set mode bit)

Cannot change address space

Cannot disable interrupts

Cannot perform IO operations

Cannot halt the processor

# How can an application do anything useful ...

---

Asks for permission to access kernel mode!

**System calls** Transition from user to kernel mode only at specific locations specified by the OS

**Exceptions** User mode code attempts to execute a privileged operation. Generates a processor exception which passes control to kernel at specific locations

More on safe control transfers later



# Hardware must support

---

## 1) Privileged Instructions

Unsafe instructions cannot be executed in user mode

## 2) Memory Isolation

Memory accesses outside a process's address space prohibited

## 3) Interrupts

Ensure kernel can regain control from running process

## 4) Safe Transfers

Correctly transfer control from user-mode to kernel-mode and back

## Req 2/4: Memory Protection

---

OS and applications both resident in memory

Application should not read/write kernel memory  
(or other apps memory)

# A Bug's Tail

---

The character could leave the game area and start overwriting other running programs and kernel memory.

One of the worst bugs I ever had to deal with was in this game. Once the game player made it to the Colony, every so often the system would crash and burn at totally random times. You might be playing for ten minutes when it happened or ten hours, but it would just die in a totally random way

There was a slow-moving slug like creature that knew how to follow the game player's trail. When it came across another creature, rather than bouncing off and risk losing the trail, I made it so that it would destroy the other creature and stay on target to find you. This worked great, except that on some rare occasions, this slug could do to a wall what it did to the other creatures. That is, it could delete it. This meant that the virtual door was now open for this creature to explore the rest of the RAM on the Macintosh, deleting and modifying it as it went along. Of course, it was just a matter of time before it found some juicy code. In other words, the bug was a REAL bug.

# Super Mario Land 2

---

Mario could exit a level and explore the entire memory of the system



# Virtual Memory is Hard!

---

Virtualizing the CPU

Process Abstraction and API

Threads and Concurrency

Scheduling

Virtualizing Memory

Virtual Memory

Paging

Persistence

IO devices

File Systems

Distributed Systems

Challenges with distribution

Data Processing & Storage

# Hardware must support

---

## 1) Privileged Instructions

Unsafe instructions cannot be executed in user mode

## 2) Memory Isolation

Memory accesses outside a process's address space prohibited

## 3) Interrupts

Ensure kernel can regain control from running process

## 4) Safe Transfers

Correctly transfer control from user-mode to kernel-mode and back

## Req 3/4: Interrupts

---

Kernel must be able to **regain control** of the processor

Hardware to the rescue! (Again x 2)

### Hardware Interrupts

Set to interrupt processor after a specified delay or specified event and transfer control to (specific locations) in Kernel.

Resetting timer is a privileged operation